

The Law of AI: Mock Exam

Vrije Universiteit Amsterdam · 2026

Student copy

Instructions

This mock exam has 30 multiple-choice questions across two parts. Part A (Q1-Q27) covers the lectures and mandatory readings. Part B (Q28-Q30) asks you to choose the best rewrite of a provision: a good rewrite is clearer than the original while remaining proportional to the provision's objective and consistent with the rest of the regulation.

Each question has exactly one correct answer. Choose the single best option.

Part A — Lectures & readings (Q1-Q27)

1. Lessig's Pathetic Dot Theory holds that the behaviour of the regulated “dot” is shaped by four modalities of regulation. Which set lists them correctly?

- A) Law, ethics, politics, and technology
- B) Legislation, case law, contracts, and standards
- C) Law, social norms, the market, and architecture (code)
- D) Law, enforcement, deterrence, and sanctions

2. In Schrepel's “Not-So-Pathetic Dot Theory” (2022), what is the central refinement of Lessig's model?

- A) Law should be removed as a constraint because code now does its work
- B) Architecture should always be treated as the most powerful of the four constraints
- C) The dot is not passive: it reacts to and reshapes the four constraints, forming a complex adaptive system
- D) The four constraints operate independently and never influence one another

3. Instagram automatically down-ranks posts its classifier flags as likely false, so fewer users see them. In Lessig's framework this intervention is best classified as regulation through:

- A) Law
- B) Architecture (code)
- C) Social norms
- D) The market

4. In Lecture 1, Schrepel drew lessons for regulators from the Not-So-Pathetic Dot Theory. Which best reflects his caution about how the law affects technology (architecture)?

- A) The law should be careful when “freezing” a technology and stopping its evolution — e.g. mandating a single fixed standard such as USB-C
- B) The law should never interact with architecture, only with norms and the market
- C) The law should always pick a single winning technology as early as possible
- D) The law should be written once and then left unchanged in order to be “future-proof”

5. Article 3(1) of the AI Act defines an “AI system.” Which feature is part of that definition?

- A) It must rely exclusively on deep neural networks

- B) It must be trained on personal data
- C) It infers, from the input it receives, how to generate outputs such as predictions, content, recommendations or decisions
- D) It must operate fully autonomously, with no human involvement

6. The AI Act's risk-based approach sorts AI into four tiers. Which set lists them correctly?

- A) Systemic, high, medium, and low risk
- B) Critical, serious, moderate, and negligible risk
- C) Unacceptable, high, limited, and minimal risk
- D) Prohibited, regulated, voluntary, and exempt

7. Which practice is prohibited outright under Article 5?

- A) An AI medical device that assists diagnosis
- B) A customer-service chatbot on a retailer's website
- C) An AI system used to screen job applications
- D) Social scoring that leads to detrimental or disproportionate treatment of people

8. The lecture noted recent additions to the Article 5 list of prohibited practices. Which of the following was newly added?

- A) Recommender systems on social-media platforms
- B) AI used to draft marketing copy
- C) AI used for email spam filtering
- D) AI that generates non-consensual intimate imagery ("nudifier" apps) and AI-generated child sexual abuse material

9. Non-compliance with the Article 5 prohibition on unacceptable-risk AI can attract fines of up to:

- A) €750,000 in all cases
- B) €10 million or 2% of total worldwide annual turnover, whichever is higher
- C) €15 million or 3% of total worldwide annual turnover, whichever is higher
- D) €35 million or 7% of total worldwide annual turnover, whichever is higher

10. Non-compliance with the obligations for high-risk AI systems (e.g. Articles 9–15) can attract fines under Article 99(4) of up to:

- A) No fine: high-risk breaches attract only corrective orders
- B) €35 million or 7% of total worldwide annual turnover whichever is higher
- C) €15 million or 3% of total worldwide annual turnover, whichever is higher
- D) €1,500,000 in all cases

11. Even where a system falls within an Annex III area, Article 6(2) provides that it is NOT high-risk if:

- A) It was developed outside the EU
- B) Its provider is a small or medium-sized enterprise
- C) It does not pose a significant risk of harm to health, safety or fundamental rights (e.g. it performs a narrow procedural task)
- D) It is offered under an open-source licence

12. Article 9 requires a risk-management system for high-risk AI systems. The Act describes this system as:

- A) A one-off assessment carried out once, before market placement
- B) A purely voluntary code of conduct
- C) An obligation that applies only to GPAI providers
- D) A continuous, iterative process maintained and updated throughout the system's entire lifecycle

13. In “Decoding the AI Act,” Schrepel discusses Article 10's requirement that datasets be “relevant, sufficiently representative and, to the best extent possible, free of errors and complete.” His assessment is that this wording:

- A) Is fully precise and leaves no room for interpretation
- B) Should be deleted because data quality is irrelevant to risk
- C) Improves on the original “free of errors” draft (which was unrealistic) but remains vague and likely to generate litigation
- D) Requires datasets to be entirely free of errors, as data scientists recommend

14. Article 11 requires technical documentation for high-risk AI systems. What accommodation does the Act make for SMEs, and how does Schrepel characterise it?

- A) SMEs face stricter documentation duties than large firms
- B) SMEs may provide the documentation in a simplified form; he sees it as the only SME-specific relief in the high-risk requirements and likely insufficient to prevent market distortion
- C) SMEs are fully exempt from all high-risk obligations
- D) The Act draws no distinction between SMEs and large firms anywhere

15. Article 15 of the AI Act requires that high-risk AI systems be designed to achieve and maintain, throughout their lifecycle, an appropriate level of:

- A) Open-source licensing and public release of the model weights
- B) Accuracy, robustness and cybersecurity
- C) Energy efficiency certified by a notified body
- D) Profitability and market share

16. “Limited risk” AI systems under Article 50 are subject mainly to:

- A) Transparency obligations — e.g. telling people they are interacting with an AI, and labelling deepfakes/synthetic content
- B) Full conformity assessment and CE marking
- C) The systemic-risk evaluation regime
- D) A complete ban

17. Which best captures the Article 3(63) definition of a “general-purpose AI model” (GPAI)?

- A) Any AI system classified as high-risk under Annex III
- B) A model that displays significant generality, can competently perform a wide range of distinct tasks and be integrated into many downstream systems (excluding models used only for R&D/prototyping before market)
- C) Any AI system deployed by a very large online platform
- D) A model trained exclusively on personal data

18. A GPAI model is presumed to have “high-impact capabilities,” and thus to pose systemic risk under Article 51, when:

- A) It has more than one billion users
- B) The cumulative compute used for training, measured in floating-point operations, is greater than 10^{25} FLOP
- C) Its provider's annual turnover exceeds €35 million
- D) It has been trained for longer than six months

19. According to “Decoding the AI Act” (citing Epoch AI, as of April 2024), which models had exceeded the systemic-risk compute threshold?

- A) Gemini Ultra, Mistral Large, GPT-4 and Inflection 2
- B) Only GPT-4
- C) No model had yet crossed the threshold
- D) Llama 2 and BLOOM

20. Article 53(1) requires providers of GPAI models (without systemic risk) to, among other things:

- A) Affix a CE marking and undergo third-party conformity assessment
- B) Put in place a policy to comply with EU copyright law — including respecting opt-outs under Art. 4(3) of the DSM Directive — and publish a sufficiently detailed summary of training content
- C) Obtain explicit consent from every person whose data appears in the training corpus
- D) Conduct a Fundamental Rights Impact Assessment under Article 27

21. Under Article 53(2), a GPAI model can benefit from reduced obligations when offered under a free and open-source licence. A key limit on this exemption is that:

- A) It does not apply to GPAI models with systemic risk
- B) It removes all obligations, including for systemic-risk models
- C) It applies only to models developed inside the EU
- D) It applies only to models trained on fewer than 10^{25} FLOP, regardless of risk

22. Schrepel & Potts (“Measuring the Openness of AI Foundation Models”) argue that the AI Act's treatment of open source is flawed because:

- A) Openness is a spectrum, yet the Act uses a binary open/closed test — risking misclassification (e.g. Llama 3 and GPT-4 differ by only 2 of 36 points)
- B) The Act gives open-source models stricter obligations than closed ones
- C) Open-source models are always more dangerous than closed ones
- D) No foundation model is open in any respect whatsoever

23. Under the revised Product Liability Directive (PLD), what is the relationship between AI Act compliance and civil liability, as explained in Lecture 5?

- A) AI Act compliance gives complete immunity from any civil liability
- B) The PLD does not apply to AI, because software is excluded from the definition of a “product”
- C) A claimant must always prove that the provider was at fault before any liability arises
- D) AI Act compliance helps rebut a defectiveness claim, while breaching mandatory safety requirements can trigger a presumption of defectiveness

24. Hupont et al. (“Documenting High-Risk AI”) assess whether existing AI transparency tools could underpin future AI Act technical standards. Which documentation approaches do they examine?

- A) Privacy notices and cookie-consent banners
- B) CE markings and customs declarations
- C) Source-code escrow agreements and patent filings
- D) Datasheets for Datasets, the Dataset Nutrition Label, Model Cards and AI FactSheets

25. A company tests a healthcare AI system inside an Article 57 regulatory sandbox. Because the sandbox operates “without prejudice to other Union law,” the company:

- A) Is exempt from the GDPR for the duration of the test
- B) May treat the sandbox authorisation as a blanket legal basis for any processing
- C) Must still identify a valid GDPR legal basis to process patient data, even for experimental purposes
- D) No longer needs to comply with the AI Act inside the sandbox

26. In “Decoding the AI Act,” Schrepel argues that the Act indirectly expands competition agencies' powers because:

- A) Market surveillance authorities may access documentation, training data and (on justified request) source code, and must pass on information of potential interest to competition agencies
- B) It empowers competition agencies to block mergers in AI markets
- C) It transfers all GPAI enforcement to national competition authorities
- D) It allows competition agencies, in all cases, to demand information without any suspicion of an infringement of competition law generally

27. The Graux et al. study for the European Parliament (ITRE) on the interplay between the AI Act and other EU digital legislation concludes that:

- A) Each instrument is individually well-targeted, but their interplay creates significant regulatory complexity, overlap and burden
- B) The AI Act should be repealed and replaced by the GDPR
- C) Only the AI Act applies to AI systems; the other instruments are displaced
- D) The frameworks already form a fully coherent system with no overlaps or gaps

Part B — Rewriting AI Act provisions (Q28-Q30)

For each provision below, choose the best rewrite.

28. REWRITE — Article 5 prohibition on manipulative AI. The current text bans AI that uses subliminal or manipulative techniques “materially distorting” behaviour and causing “significant harm.” Which rewrite is best — clearer while staying proportional to the provision’s objective (protecting autonomy) and consistent with the rest of the Act?

- A) “Prohibited: AI systems using subliminal or purposefully manipulative techniques that appreciably impair a person’s ability to make an informed decision and thereby cause, or are likely to cause, significant harm.”
- B) “Prohibited: AI systems that are manipulative.”
- C) “Prohibited: subliminal AI techniques, but only where the provider expressly admits an intention to cause harm.”
- D) “Prohibited: any AI system that influences a user’s behaviour.”

29. REWRITE — Article 10(3) data quality. The current text requires datasets to be “relevant, sufficiently representative and, to the best extent possible, free of errors and complete.” Which rewrite is best — clearer while remaining proportional to the objective (data quality / non-discrimination) and realistic for the rest of the regime?

- A) “Datasets shall be completely free of all errors.”
- B) “Datasets shall be appropriate.”
- C) “Datasets may contain whatever data the provider chooses.”
- D) “Training, validation and testing datasets shall be relevant and sufficiently representative of the persons and contexts of the intended purpose, and subject to reasonable, documented measures to detect and reduce errors.”

30. REWRITE — Article 14 human oversight. The current text requires that persons assigned oversight can “remain aware” of the tendency to over-rely on outputs (automation bias). Which rewrite is best — clearer while staying proportional to the objective (effective human control) and workable across high-risk systems?

- A) “Providers shall enable deployers to assign oversight to persons who have the competence, authority and means to monitor the system, correctly interpret its output, and decide not to use it or to override its output.”
- B) “Persons overseeing the system should stay aware of its risks.”
- C) “Human oversight is recommended where feasible.”
- D) “A human must individually approve every single output before it takes effect.”